

Here is the Agile implementation plan for the backend. I have broken the work into **3 Sprints**, with small, isolated tasks suitable for individual branches.

Sprint 1: Foundation & Data Ingestion (The Producer)

Goal: Build the Core domain and the "Write" side of the system (Downloading, Parsing, and Saving data).

Task ID	Branch Name	Description & Files	Testing Strategy
1.1	feature/core-domain	Implement Domain Entities & Enums Create the foundational data structures. Files: Dataset.cs, MetadataRecord.cs, MetadataFormat.cs, Result.cs.	Unit Tests: Verify that entity constructors validation works (e.g., ensuring a Dataset cannot exist without an ID).
1.2	feature/infra-sqlite	Implement SQLite Repository Set up Entity Framework Core and the database schema. Files: AppDbContext.cs, SqliteMetadataRepository.cs.	Integration Tests: Use SqliteMetadataRepositoryTests.cs with an In-Memory SQLite connection to ensure CRUD operations work without a real disk file.

1.3	feature/parsing-strategies	<p>Implement Metadata Strategies</p> <p>Create the parsers for XML, JSON-LD, etc.</p> <p>Files: MetadataParserFactory.cs, Iso19115XmlParser.cs, JsonExpandedParser.cs.</p>	<p>Unit Tests: Create ParserTests.cs. Feed raw sample XML/JSON strings into each parser and assert that the output MetadataRecord matches expected values.</p>
1.4	feature/download-extract	<p>Implement Download & Zip Extraction</p> <p>Handle HTTP streams and Zip file inspection.</p> <p>Files: CehDatasetDownloader.cs, ZipExtractionService.cs.</p>	<p>Integration Tests: Test ZipExtractionService by creating a small dummy .zip in the test setup, passing it to the service, and verifying it identifies the correct files inside.</p>
1.5	feature/etl-orchestrator	<p>Implement ETL Workflow</p> <p>Wire everything together: Download -> Unzip -> Parse -> Save.</p> <p>Files: EtlOrchestrator.cs, DatasetProcessor.cs.</p>	<p>Unit Tests: Mock all interfaces (IDownloader, IRepository). Test the <i>flow</i>: "If download succeeds, does it call Parse? If Parse succeeds, does it call Save?"</p>

Sprint 2: Search & Intelligence (The Consumer)

Goal: Build the "Read" side of the system (Vector Storage, Search Logic, and RAG).

Task ID	Branch Name	Description & Files	Testing Strategy
2.1	feature/infra-vector	Implement Vector Store Adapter Connect to the Qdrant instance. Files: QdrantAdapter.cs, EmbeddingVector.cs.	Integration Tests: Use the Docker Qdrant container. Index a dummy vector and try to retrieve it via SearchAsync.
2.2	feature/search-service	Implement Search Logic Combine keyword search (SQLite) with semantic search (Qdrant). Files: SearchService.cs, SearchResultDto.cs.	Unit Tests: Mock the Repositories. Test logic like "If vector store returns ID 5, does the service fetch Metadata for ID 5 from SQLite?"
2.3	feature/rag-agent	Implement Conversational Agent	Unit Tests: Mock the ILLmClient. Verify that RagService constructs the prompt string correctly (e.g., "Context:

		<p>Build the prompt engineering logic.</p> <p>Files: RagService.cs, OpenAiLlmClient.cs.</p>	[Data]... Question: [User Query]").
--	--	--	-------------------------------------

Sprint 3: API & Delivery

Goal: Expose functionality to the web and harden the system.

Task ID	Branch Name	Description & Files	Testing Strategy
3.1	feature/api-controllers	<p>Implement Web API Endpoints</p> <p>Connect HTTP requests to the Application Core.</p> <p>Files: SearchController.cs, ChatController.cs, AdminEtlController.cs.</p>	<p>End-to-End Tests: Use WebApplicationFactory. Spin up the full API in memory and send a real HTTP GET request to /api/search to ensure the full pipeline works.</p>
3.2	feature/api-caching	<p>Implement Response Caching</p> <p>Add middleware to cache frequent search results.</p> <p>Files: ResponseCachingMiddleware.cs.</p>	<p>Integration Tests: Send the same request twice. Verify the second request returns faster or has a specific "Cached" header.</p>

3.3	chore/documentation	<p>Swagger & Readme</p> <p>Finalize API documentation.</p> <p>Files: Program.cs (Swagger Config), README.md.</p>	<p>Manual Verification:</p> <p>Run the container and verify the Swagger UI is accessible.</p>
------------	---------------------	--	--

Suggested First Step

I recommend we start with **Task 1.1 (Domain Foundation)**.

Command to start:

Bash

```
git checkout -b feature/core-domain
```

Would you like me to generate the **C# code for Task 1.1** (The Entity classes and Enums) so you can push the first branch?